



OpenBioML PyTorch Lightning workshop (1/2)

Adrian Wälchli, Research Engineer
Luca Antiga, CTO

OpenBioML PyTorch Lightning workshop

Session 1: Thu 23 Feb, 3pm ET

Session 2: Thu 2 Mar, 3pm ET

<https://harvard.zoom.us/j/97375262666>

OpenBioML PyTorch Lightning workshop

Session 1:

- Intro to PyTorch Lightning + Fabric
- Hands-on: raw PyTorch -> Fabric -> PyTorch Lightning Trainer
- A look into OpenFold

By the end of Session 1 you will know how to build a model with Lightning and train it. Distributed. On a SLURM cluster.

OpenBioML PyTorch Lightning workshop

Session 2:

- Intro to core distributed concepts
- Hands-on: how to debug and optimize performance, single node and distributed, running benchmarks
- More on OpenFold
- Quick look at the Lightning platform

By the end of Session 2 you will know how to make sure you are setting up your training correctly and verify you are leveraging your hardware the best.

Join us here

discord

discord.gg/MWAEvnC5fU

forums

lightning.ai/forums

twitter

[@LightningAI](https://twitter.com/LightningAI)



<https://linktr.ee/lightningai>

Intro: PyTorch Lightning and Lightning Fabric

PyTorch Lightning in the wild

Stable diffusion




<https://github.com/Stability-AI/stablediffusion>

<https://github.com/CompVis/stable-diffusion>

PyTorch Lightning in the wild


Instruct Pix2Pix



<https://github.com/timothybrooks/instruct-pix2pix>

PyTorch Lightning in the wild

OpenFold



<https://github.com/aqlaboratory/openfold>

<https://www.biorxiv.org/content/10.1101/2022.11.20.517210v1.full.pdf>

PyTorch Lightning in the wild

NeMo Megatron (framework)



part of NeMo

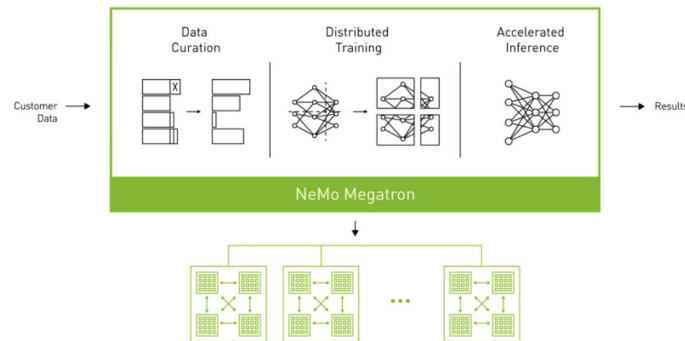
<https://nvidia.github.io/NeMo/>

Acoustic
Decoders
Language
BERT
Speech Synthesis
Vision Encoder

<https://developer.nvidia.com/nemo/megatron>

NVIDIA NeMo Megatron is an end-to-end framework for training and deploying LLMs with billions and trillions of parameters.

[Download Now](#)



PyTorch Lightning

You do the science, we do the engineering

33 million+
DOWNLOADS

780+
CONTRIBUTORS

 pip install lightning 21,512

13,000+
PROJECTS USING LIGHTNING

6,000+
SLACK MEMBERS 

10,000+ ORGANIZATIONS BUILD WITH LIGHTNING



<https://lightning.ai>

PyTorch Lightning

Organized PyTorch

LightningModule

```
import lightning as L
from torch import nn, optim

encoder = nn.Sequential(nn.Linear(28 * 28, 64), nn.ReLU(), nn.Linear(64, 3))
decoder = nn.Sequential(nn.Linear(3, 64), nn.ReLU(), nn.Linear(64, 28 * 28))
```

```
class LitAutoEncoder(L.LightningModule):
    def __init__(self, encoder, decoder):
        super().__init__()
        self.encoder = encoder
        self.decoder = decoder

    def training_step(self, batch, batch_idx):
        X, y = batch
        X = x.view(x.size(0), -1)
        z = self.encoder(x)
        x_hat = self.decoder(z)
        loss = nn.functional.mse_loss(x_hat, x)
        self.log("train_loss", loss) return loss

    def configure_optimizers(self):
        optimizer = optim.Adam(self.parameters(), lr=1e-3)
        return optimizer
```

```
autoencoder = LitAutoEncoder(encoder, decoder)
dataset = MIST(os.getcwd(), download=True, transform=ToTensor())

train_loader = utils.data.DataLoader (dataset)
```

```
trainer = L.Trainer(limit_train_batches=100, max_epochs=1)
trainer.fit(model=autoencoder, train_dataloaders=train_loader)
```

Trainer

PyTorch Lightning

Accelerators

GPU, TPU, HPU, IPU, MPS

Strategies

DDP, FSDP, DeepSpeed, Colossal AI

Precision

Callbacks

```
# train on 4 GPUs
trainer = Trainer(
    devices=4,
    accelerator="gpu"
)

# train 1B+ parameter models with Deepspeed/fsdp
trainer = Trainer(
    devices=4,
    accelerator="gpu",
    strategy="deepspeed_stage_2",
    precision=16
)

# 20+ helpful flags for rapid idea iteration
trainer = Trainer(
    max_epochs=10,
    min_epochs=5,
    overfit_batches=1
)

# access the latest state of the art techniques
trainer = Trainer(callbacks=[StochasticWeightAveraging(...)])
```

PyTorch Lightning

The anatomy of `.fit()`

```
def fit(self):
    if global_rank == 0:
        # prepare data is called on GLOBAL ZERO only
        prepare_data()

    configure_callbacks()

    with parallel(devices):
        # devices can be GPUs, TPUs,
        train_on_device(model)

def train_on_device(model):
    # called PER DEVICE
    setup("fit")
    configure_optimizers()
    on_fit_start()

    # sanity check runs here
    on_train_start()

    for epoch in epochs:
        fit_loop()

    on_train_end()
    on_fit_end()
    teardown("fit")
```

```
def fit_loop():
    on_train_epoch_start()

    for batch in train_dataloader():
        on_train_batch_start()

        on_before_batch_transfer()
        transfer_batch_to_device()
        on_after_batch_transfer()

        training_step()

        on_before_zero_grad()
        optimizer_zero_grad()

        on_before_backward()
        backward()
        on_after_backward()

        on_before_optimizer_step()
        configure_gradient_clipping()
        optimizer_step()

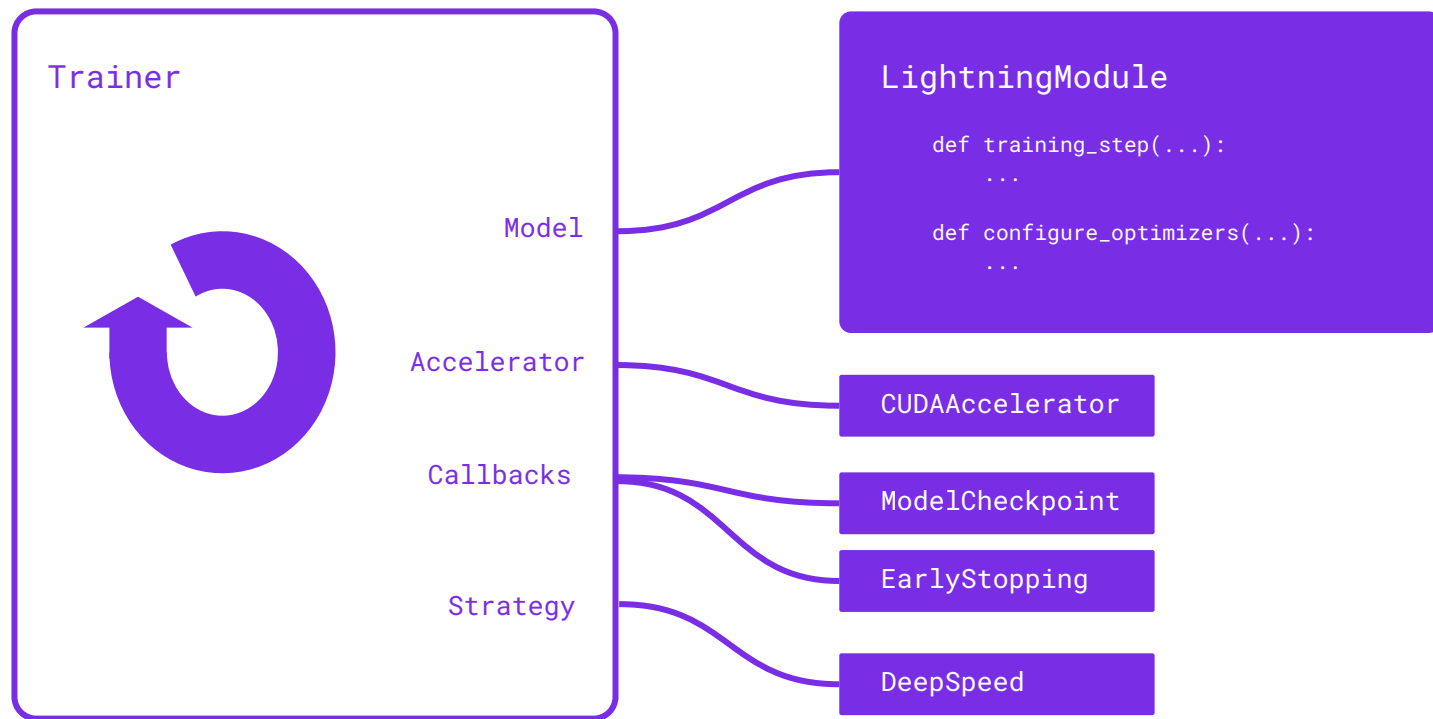
    on_train_batch_end()

    if should_check_val:
        val_loop()

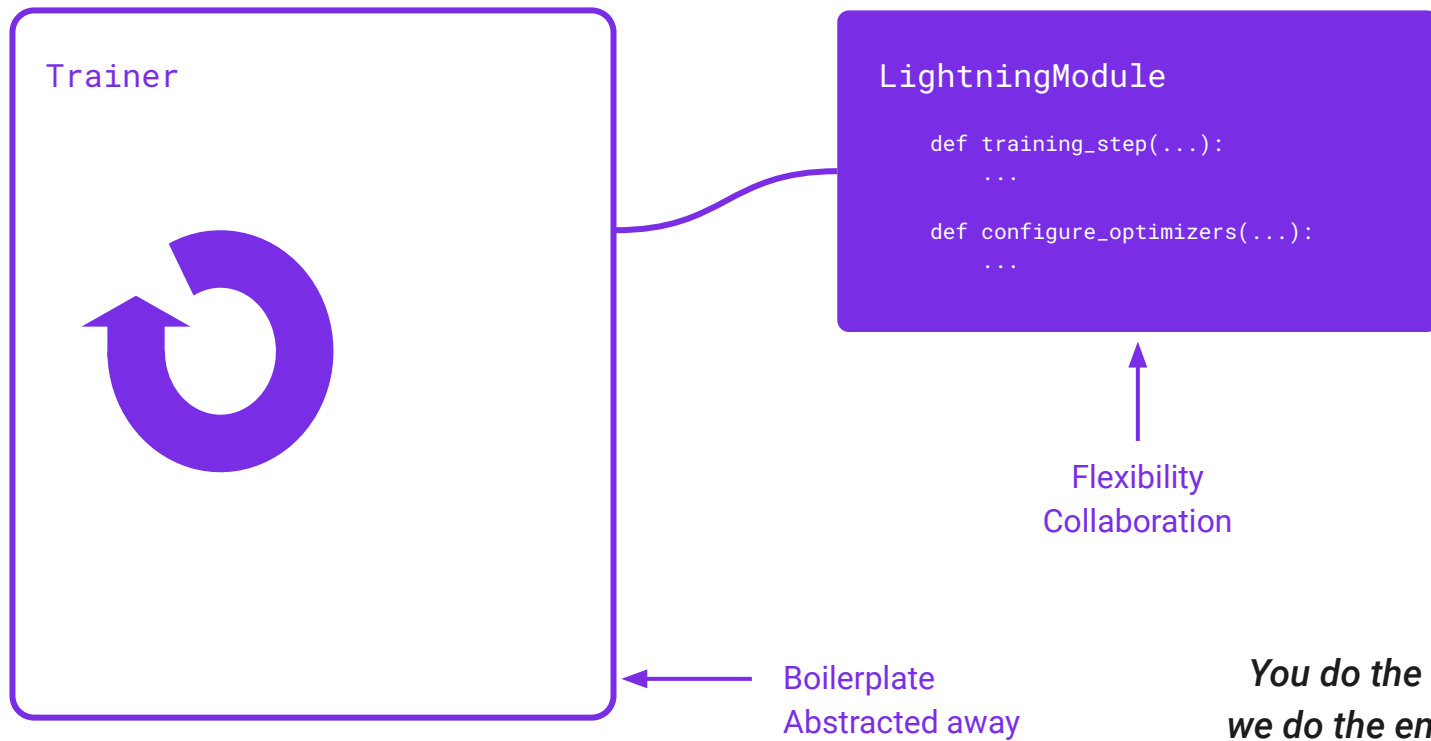
    # end training epoch
    training_epoch_end()

    on_train_epoch_end()
```

PyTorch Lightning



PyTorch Lightning



Get Started

Lightning in 15 minutes

Installation

Level Up

Basic skills

Intermediate skills

Advanced skills

Expert skills

Core API

LightningModule

Trainer

Fabric (Beta)

API Reference

accelerators

callbacks

cli

core

loggers

profiler

trainer

strategies

pytorch-lightning.readthedocs.io

Docs > Train 1 trillion+ parameter models

Edit on GitHub

Shortcuts

TRAIN 1 TRILLION+ PARAMETER MODELS

When training large models, fitting larger batch sizes, or trying to increase throughput using multi-GPU compute, Lightning provides advanced optimized distributed training strategies to support these cases and offer substantial improvements in memory usage.

Note that some of the extreme memory saving configurations will affect the speed of training. This Speed/Memory trade-off in most cases can be adjusted.

Some of these memory-efficient strategies rely on offloading onto other forms of memory, such as CPU RAM or NVMe. This means you can even see memory benefits on a **single GPU**, using a strategy such as [DeepSpeed ZeRO Stage 3 Offload](#).

Check out this amazing video explaining model parallelism and how it works behind the scenes:

NVIDIA GTC 21: Half The Memory with Zero Co...
DeepSpeed Training

Watch on YouTube

Train 1 trillion+ parameter models

+ Choosing an Advanced Distributed GPU Strategy

- Colossal-AI

Placement Policy

Sharded Training

- Fully Sharded Training

Auto Wrapping

Manual Wrapping

Activation Checkpointing

- DeepSpeed

DeepSpeed ZeRO Stage 1

+ DeepSpeed ZeRO Stage 2

+ DeepSpeed ZeRO Stage 3

Custom DeepSpeed Config

+ DDP Optimizations

17

Choosing an Advanced Distributed GPU Strategy

Lightning^{AI}

Training in the age of foundation models


Sometimes the science *is* the engineering

Large models, long projects


Often very little iteration on the models, but
need of full control on the training process



A binary choice until today



Introducing Fabric



Fabric

3 Your code

```
import lightning as L

def train(fabric, model, optimizer, dataloader):
    # Training loop
    model.train()
    for epoch in range(num_epochs):
        for i, batch in enumerate(dataloader):
            ...
```

1 Fabric object

```
# Configure Fabric
fabric = L.Fabric(...)
```

```
# Instantiate objects
model = ...
optimizer = ...
train_dataloader = ...
```

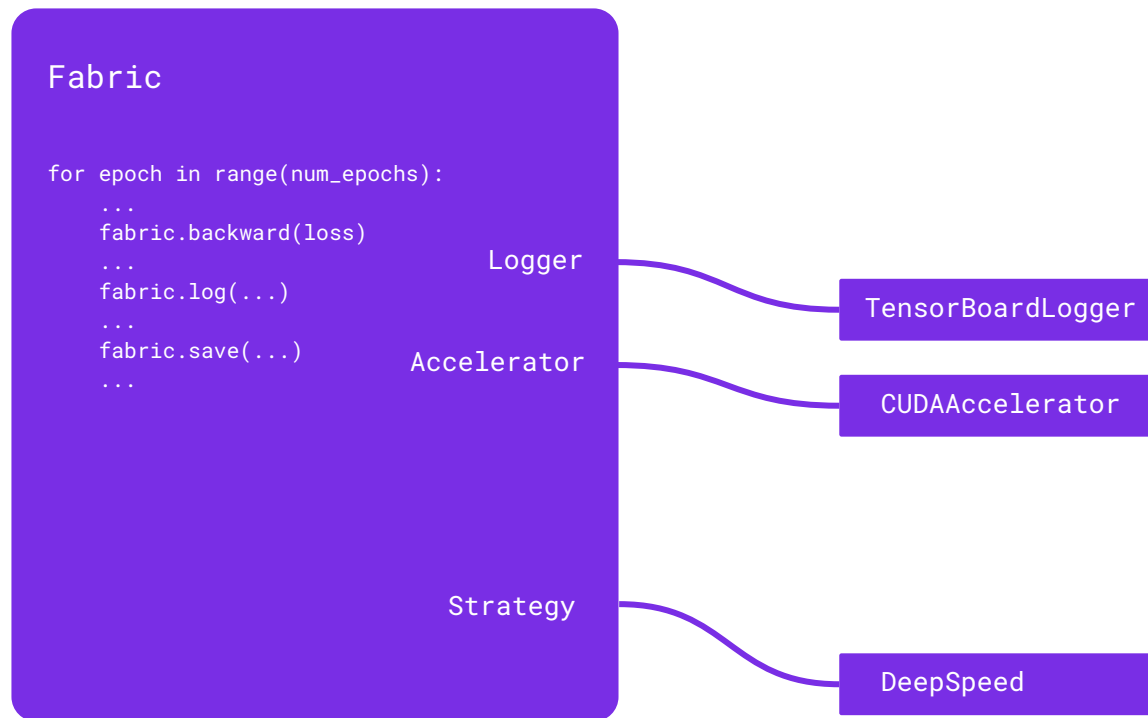
2 Setup model, optimizer,
dataloader

```
# Set up objects
model, optimizer = fabric.setup(model, optimizer)
train_dataloader = fabric.setup_dataloaders(train_dataloader)
```

```
# Run training loop
train(fabric, model, optimizer, train_dataloader)
```

```
if __name__ == "__main__":
    main()
```

Fabric



Fabric

Accelerators

Strategies

Precision

```
fabric = Fabric( devices=4, accelerator="gpu", strategy="ddp", precision=16)
```

Fabric

Loggers

```
from lightning.fabric import Fabric
from lightning.fabric.loggers import CSVLogger,
TensorBoardLogger

tb_logger = TensorBoardLogger(root_dir="logs/tensorboard")
csv_logger = CSVLogger(root_dir="logs/csv")

# Add multiple loggers in a list
fabric = Fabric(loggers=[tb_logger, csv_logger])

# Calling .log() or .log_dict() always logs to all loggers
simultaneously
fabric.log("some_value", value)
```


Fabric

Checkpoints


```
# Define the state of your program/loop
state = {"model1": model1, "model2": model2, "optimizer": optimizer,
        "iteration": iteration, "hparams": ...}

# Save a checkpoint
fabric.save("path/to/checkpoint ckpt", state)


# Load a checkpoint
fabric.load("path/to/checkpoint.ckpt", state)

# Restore part of a checkpoint
state = {"model1": model1}
remainder = fabric.load("path/to/checkpoint.ckpt", state)
```

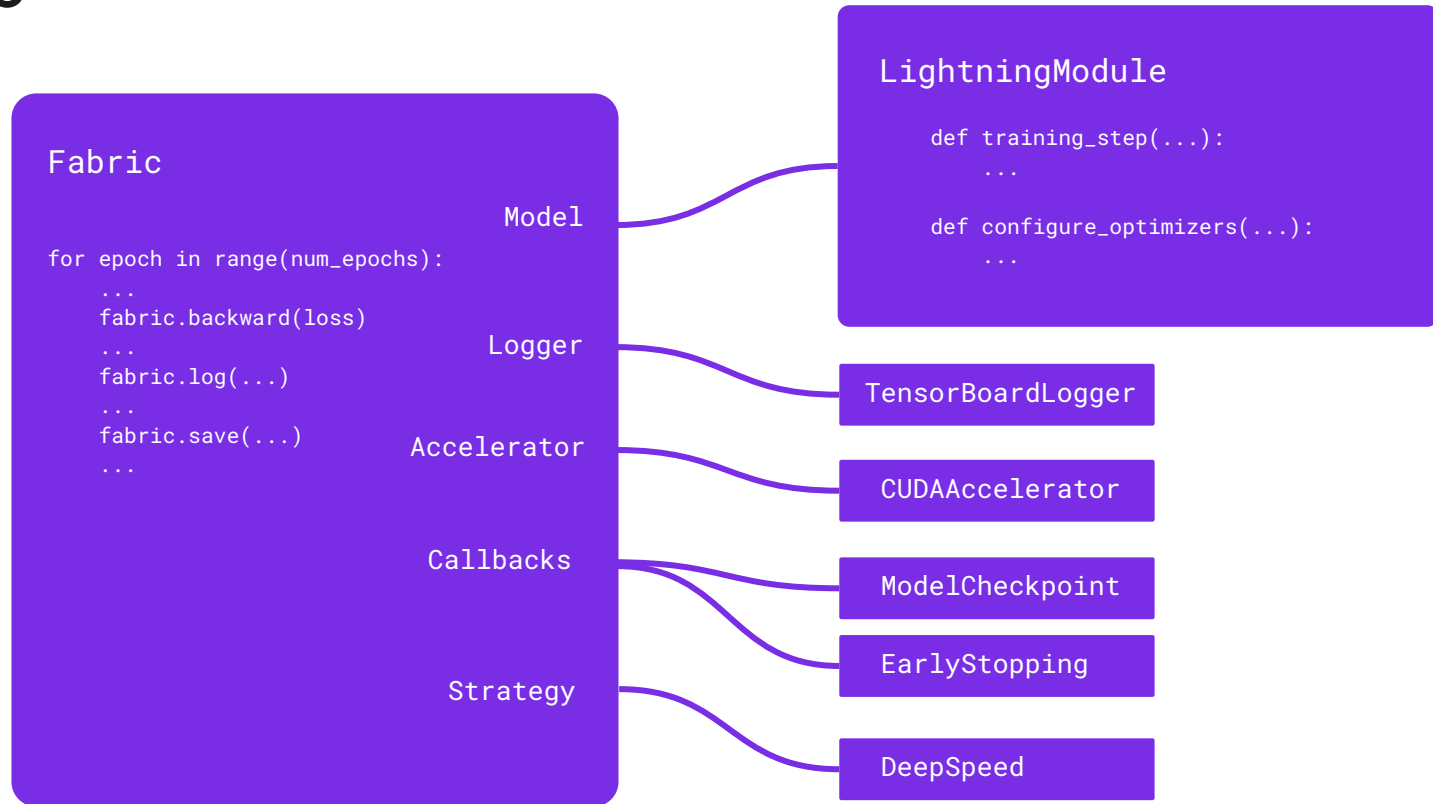
Fabric



Fabric



Fabric



Fabric

LightningModule

```
import lightning as L

class LitModel(L.LightningModule):
    def __init__(self):
        super().__init__()
        self.model = ...

    def training_step(self, batch, batch_idx):
        # Main forward, loss computation, and metrics goes here
        x, y = batch
        y_hat = self.model(x)
        loss = self.loss_fn(y, y_hat)
        acc = self.accuracy(y, y_hat)
        ...
        return loss

    def configure_optimizers(self):
        # Return one or several optimizers
        return torch.optim.Adam(self.parameters() , ...)

    def train_dataloader(self):
        # Return your dataloader for training
        return DataLoader(...)

    def on_train_start(self):
        # Do something at the beginning of training
        ...

    def any_hook_you_like(self, *args, **kwargs):
        ...
```

Fabric

LightningModule

Training code is
model-independent

```
import lightning as L

fabric = L.Fabric(...)

# Instantiate the LightningModule
model = LitModel()

# Get the optimizer(s) from the LightningModule
optimizer = model.configure_optimizers()

# Get the training data loader from the LightningModule
train_dataloader = model.train_dataloader()


# Set up objects
model, optimizer = fabric.setup(model, optimizer)
train_dataloader = fabric.setup_dataloaders(train_dataloader)

# Call the hooks at the right time
model.on_train_start()

model.train()
for epoch in range(num_epochs) :
    for i, batch in enumerate(dataloader):
        optimizer.zero_grad()
        loss = model.training_step(batch, i)
        fabric.backward(loss)
        optimizer.step()

# Control when hooks are called
if condition:
    model.any_hook_you_like()
```

Fabric



Fabric

Callbacks

```
from lightning.fabric import Fabric

# The code of a callback can live anywhere, away from the training loop
from my_callbacks import MyCallback

# Add one or several callbacks:
fabric = Fabric(callbacks=[MyCallback()])

for iteration, batch in enumerate(train_dataloader):
    ...
    fabric.backward(loss)
    optimizer.step()

    # Let a callback add some arbitrary processing at the appropriate
    place
    # Give the callback access to some variables
    fabric.call("on_train_batch_end", loss=loss, output=...)
```


Fabric

Callbacks

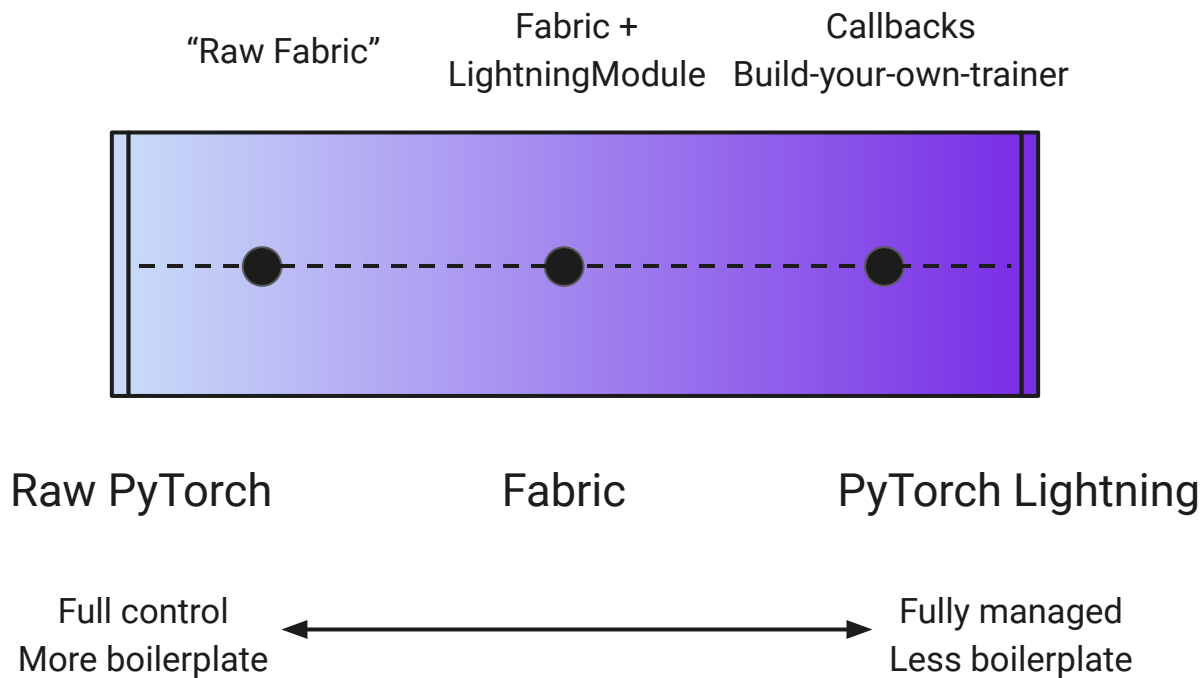
```
# Add multiple callback implementations in a list
callback1 = LearningRateMonitor()
callback2 = Profiler()

fabric = Fabric(callbacks=[callback1, callback2])

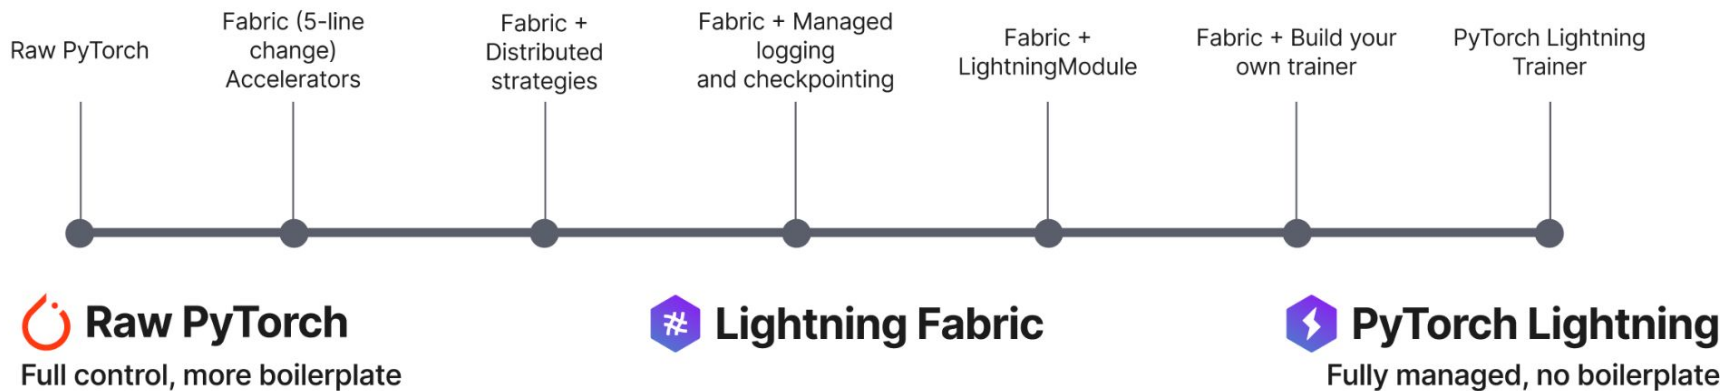
# Let Fabric call the implementations (if they exist)
fabric.call("any_callback_method", arg1=..., arg2=...)

# fabric.call is the same as doing this
callback1.any_callback_method(arg1=..., arg2=...)
callback2.any_callback_method(arg1=..., arg2=...)
```

Fabric



Fabric



Communication

boilerplate




```
116 # Build the data-parallel groups.
117 global _DATA_PARALLEL_GROUP
118 global _DATA_PARALLEL_GLOBAL_RANKS
119 assert _DATA_PARALLEL_GROUP is None, 'data parallel group is already initialized'
120 all_data_parallel_group_ranks = []
121 for i in range(pipeline_model_parallel_size):
122     start_rank = i * num_pipeline_model_parallel_groups
123     end_rank = (i + 1) * num_pipeline_model_parallel_groups
124     for j in range(tensor_model_parallel_size):
125         ranks = range(start_rank + j, end_rank, tensor_model_parallel_size)
126         all_data_parallel_group_ranks.append(list(ranks))
127         group = torch.distributed.new_group(ranks)
128         if rank in ranks:
129             _DATA_PARALLEL_GROUP = group
130             _DATA_PARALLEL_GLOBAL_RANKS = ranks
131
132 # Build the model-parallel groups.
133 global _MODEL_PARALLEL_GROUP
134 assert _MODEL_PARALLEL_GROUP is None, 'model parallel group is already initialized'
135 for i in range(data_parallel_size):
136     ranks = [data_parallel_group_ranks[i]
137              for data_parallel_group_ranks in all_data_parallel_group_ranks]
138     group = torch.distributed.new_group(ranks)
139     if rank in ranks:
140         _MODEL_PARALLEL_GROUP = group
141
142 # Build the tensor model-parallel groups.
143 global _TENSOR_MODEL_PARALLEL_GROUP
144 assert _TENSOR_MODEL_PARALLEL_GROUP is None, \
145     'tensor model parallel group is already initialized'
146 for i in range(num_tensor_model_parallel_groups):
147     ranks = range(i * tensor_model_parallel_size,
148                  (i + 1) * tensor_model_parallel_size)
149     group = torch.distributed.new_group(ranks)
150     if rank in ranks:
151         _TENSOR_MODEL_PARALLEL_GROUP = group
152
153 # Build the pipeline model-parallel groups and embedding groups
154 # (first and last rank in each pipeline model-parallel group).
155 global _PIPELINE_MODEL_PARALLEL_GROUP
```

from Megatron-LM

forked in BLOOM, GPT-NeoX

<https://github.com/NVIDIA/Megatron-LM>

Fabric communication API



```
from lightning.fabric import Fabric
```

```
# Devices and num nodes determine how many processes there are  
fabric = Fabric(devices=2, num_nodes=3)  
fabric.launch()
```

```
# The total number of processes running across all devices and nodes  
fabric.world_size # 2 * 3 = 6
```

```
# The global index of the current process across all devices and nodes  
fabric.global_rank # -> {0, 1, 2, 3, 4, 5}
```


```
# The index of the current process among the processes running on the  
local node  
fabric.local_rank # -> {0, 1}
```

```
# The index of the current node  
fabric.node_rank # -> {0, 1, 2}
```

```
# Do something only on rank 0  
if fabric.global_rank == 0:  
    ...
```

Fabric communication API

Barrier




```
if fabric.global_rank == 0:
    print("Downloading dataset. This can take a while ...")
    download_dataset()

# All other processes wait here until rank 0 is done with
# downloading:
fabric.barrier()

# After everyone reached the barrier, they can access the
# downloaded files:
load_dataset()
```

Fabric communication API

Broadcast



```
fabric = Fabric(devices=4, accelerator="gpu")  
fabric.launch()
```

```
# Data is different on each process
```

```
learning_rate = torch.rand(1)
```

```
print("Before broadcast:", learning_rate)
```


```
# Transfer the tensor from one process to all the others
```

```
learning_rate = fabric.broadcast(learning_rate)
```

```
print("After broadcast:", learning_rate)
```

Fabric communication API

Gather



```
fabric = Fabric(devices=4, accelerator="gpu")
fabric.launch()
```

```
# Data is different in each process
```

```
result = torch.tensor(10 * fabric.global_rank)
```


```
# Every process gathers the tensors from all other processes  
# and stacks the result:
```

```
result = fabric.all_gather(data)
```

```
print("Result of all-gather:", result) # tensor([0, 10, 20, 30])
```


Fabric communication API

Reduce



```
fabric = Fabric(devices=4, accelerator="gpu")
fabric.launch()

# Data is different in each process
data = torch.tensor(10 * fabric.global_rank)


# Sum the tensors from every process
result = fabric.all_reduce(data, reduce_op="sum")

# sum (0 + 10 + 20 + 30) = tensor(60)
print("Result of all-reduce:", result)
```

Lightning 2.0

Coming mid-March 2023

Lightning 2.0



* PyTorch Lightning 1.9.x will be supported long-term

Lightning 2.0

Trainer 2.0: tighter, fewer abstractions, fewer dependencies

Super-stable API for Trainer and LightningModule

More opt-in, less opt-out

Full control with Fabric

Full PyTorch 2.0 support

PyTorch Lightning (Trainer 2.0)

Fabric

App

Demo: From PyTorch to Fabric to Trainer

github.com/Lightning-AI/open-bio-ml-workshop

Data Loading

“If Lightning can iterate over it, it can load the data.”

```
for data in dataloader:  
    ...
```

Data Loading

- PyTorch DataLoader
- TorchData (DataLoader2)
- WebDataset
- NVIDIA DALI
- ...

Data Loading

Trainer accepts training-, validation-, and test data iterables

Performs **sanity checks** for potential user error: Shuffling, num workers, uneven lengths (DDP), ...

```
trainer = Trainer(...)

train_data = DataLoader(..., batch_size=16, num_workers=4)
val_data = DataLoader(..., batch_size=8, num_workers=2)
test_data = DataLoader(..., batch_size=8, num_workers=2)

trainer.fit(model, train_data, val_data)
trainer.test(model, test_data)
```


Data Loading

LightningModule organizes
data-loading code (optional)

Hooks get called by Trainer at
the right time

```
import lightning as L

class LitModel(L.LightningModule):
    def __init__(self):
        super().__init__()
        ...

    def training_step(self, batch, batch_idx):
        ...


    def train_dataloader(self):
        return DataLoader(...)

    def val_dataloader(self):
        return DataLoader(...)

    def test_dataloader(self):
        return DataLoader(...)
```

Distributed Data Loading

Data-parallel Distributed Training



```
from torch.utils.data import DistributedSampler, DataLoader

if args.distributed:
    train_sampler = DistributedSampler(train_dataset)
    val_sampler = DistributedSampler(val_dataset, shuffle=False,
                                    drop_last=True)
else:
    train_sampler = None
    val_sampler = None

train_loader = DataLoader(
    train_dataset, batch_size=args.batch_size, shuffle=(train_sampler is
    None),
    num_workers=args.workers, pin_memory=True, sampler=train_sampler
)

val_loader = DataLoader(
    val_dataset, batch_size=args.batch_size, shuffle=False,
    num_workers=args.workers, pin_memory=True, sampler=val_sampler
)
```

Warning:
Boilerplate code!

```
for epoch in range(args.start_epoch, args.epochs):
    if args.distributed:
        train_sampler.set_epoch(epoch)
```

Distributed Data Loading

Fabric

```
fabric = Fabric(...)

train_data = DataLoader(..., batch_size=16, num_workers=4)
val_data = DataLoader(..., batch_size=8, num_workers=2)
test_data = DataLoader(..., batch_size=8, num_workers=2)

train_data, val_data = fabric.setup_dataloaders(train_data, val_data)

for batch in train_data:
    ....

test_data = fabric.setup_dataloaders(test_data)
...
```

Trainer

```
trainer = Trainer(...)

train_data = DataLoader(..., batch_size=16, num_workers=4)
val_data = DataLoader(..., batch_size=8, num_workers=2)
test_data = DataLoader(..., batch_size=8, num_workers=2)

trainer.fit(model, train_data, val_data)
trainer.test(model, test_data)
```

Multi-node: SLURM

Step 1: Set devices and number of nodes

```
# Fabric
fabric = Fabric(accelerator="gpu", devices=8, num_nodes=4)
fabric.launch()
...

# OR

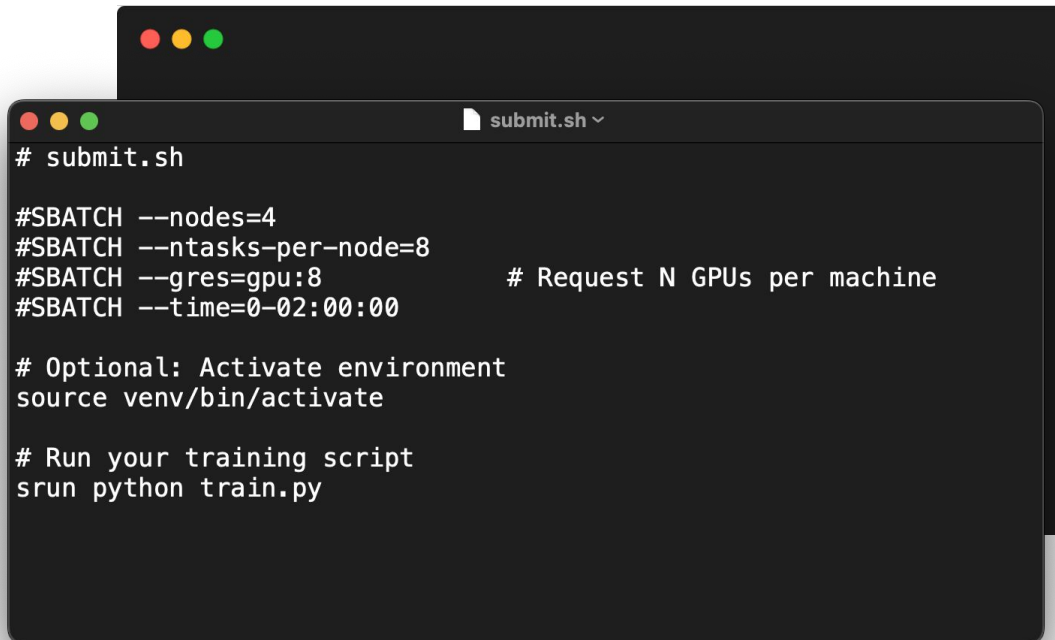
# Trainer
trainer = Trainer(accelerator="gpu", devices=8, num_nodes=4)
trainer.fit(...)
```

Multi-node: SLURM

Step 1: Set devices and number of nodes

Step 2: Create SLURM job submission script

[Template in our docs!](#)



```
# submit.sh

#SBATCH --nodes=4
#SBATCH --ntasks-per-node=8
#SBATCH --gres=gpu:8           # Request N GPUs per machine
#SBATCH --time=0-02:00:00

# Optional: Activate environment
source venv/bin/activate

# Run your training script
srun python train.py
```

Multi-node: SLURM


Step 1: Set devices and
number of nodes

Step 2: Create SLURM job
submission script

Step 3: Submit the job



```
> sbatch submit.sh
```



Walkthrough: OpenFold – AlphaFold 2 with Lightning

Reach out!

discord

discord.gg/MWAEvnC5fU

forums

lightning.ai/forums

twitter

[@LightningAI](https://twitter.com/LightningAI)



<https://linktr.ee/lightningai>

Thanks

Adrian Wälchli
Akihiro Nitta
Carlos Mocholí
Eden Afek
Ethan Harris
Jirka Borovec
Justus Schock
Thomas Chaton
William Falcon

+

♥ community



<https://linktr.ee/lightningai>

Thanks

